



School of Computing and Information Systems

INFO90002

## Database Systems and Information Modelling **SOLUTION**

Practice Exam 1

Semester 1 2021

**Reading Time: 15 minutes**

**Writing Time: 120 minutes (2 Hours)**

**This exam has 10 pages including this page  
ATTEMPT ALL QUESTIONS IN ANY ORDER**

### Authorised Materials:

**None**

### While you are undertaking this assessment, you **MUST NOT**

- make use of any messaging or communication technology
- record, screenshot, stream, upload or in any known format duplicate this document
- record, screenshot, stream, upload or in any known format duplicate your solutions
- make use of any world wide web or internet based resources such as wikipedia, github, stackoverflow, google, Weichat or any known search engine / messaging services
- act in a manner that could be regarded as providing assistance to a student who is undertaking this assessment or in the future will be undertaking this assessment
- seek assistance from any other student who is undertaking this assessment or in the future will be undertaking this assessment

### Instructions to Students

- The total for this exam is 100 marks
- Attempt **all** 9 questions which are of unequal marks value
- Be sure to number your answers to ensure your questions is marked
- Questions can be answered in any order

- **PLEASE DO NOT USE RED font colour or pens.**
- You may revisit/edit your exam answers throughout
- You must not communicate with other students whilst taking this exam, e.g. using messaging, chat rooms or email

### **IMPORTANT**

- **Your file upload must be a single Word docx document before the elapsed time.**
- **No other document format will be assessed (e.g. Pages, txt, .SQL, etc).**
- **Email submissions will not be assessed.**
- **Every question attempt must be numbered (e.g. Q2C, Q1, Q6) to ensure it is assessed.**
- **The official exam language is English. Sections of the submission in languages other than English will NOT be assessed and will be marked as 0.**
- Before submitting your solution document, check that the diagram(s) is/are readable. It is your responsibility to ensure your answers are readable and make sense to the marker.

The work you submit **must be based on your own knowledge and skills** and without the assistance of any other person. You must not copy directly the text from any material without attribution.

## Q1. ER Modelling (20 marks)

**Q1. Criterion Classics** is an on-demand streaming service that streams old television shows (e.g. 'Skippy') and movies ('Lawrence of Arabia'), including black-and-white movies ('Casablanca') and television shows ('I dream of Lucy'). Currently the company offers 3 subscriptions plans. Each plan has a code (B, S or P), name (Basic, Standard, Premium) and monthly price.

Each Criterion Classics account can have multiple profiles associated with it – for example, a family account could have one profile for the parents and one profile for each child. Profiles of type 'Child' can only view titles classified as 'U' (Universal) and 'G' (General). Each profile is stored with an avatar or digital image. About each account we store a name ('Mary Lawson'), the monthly subscription price, the number of concurrent screens allowed, the day of the month (18<sup>th</sup>) the account is charged and the subscription plan.

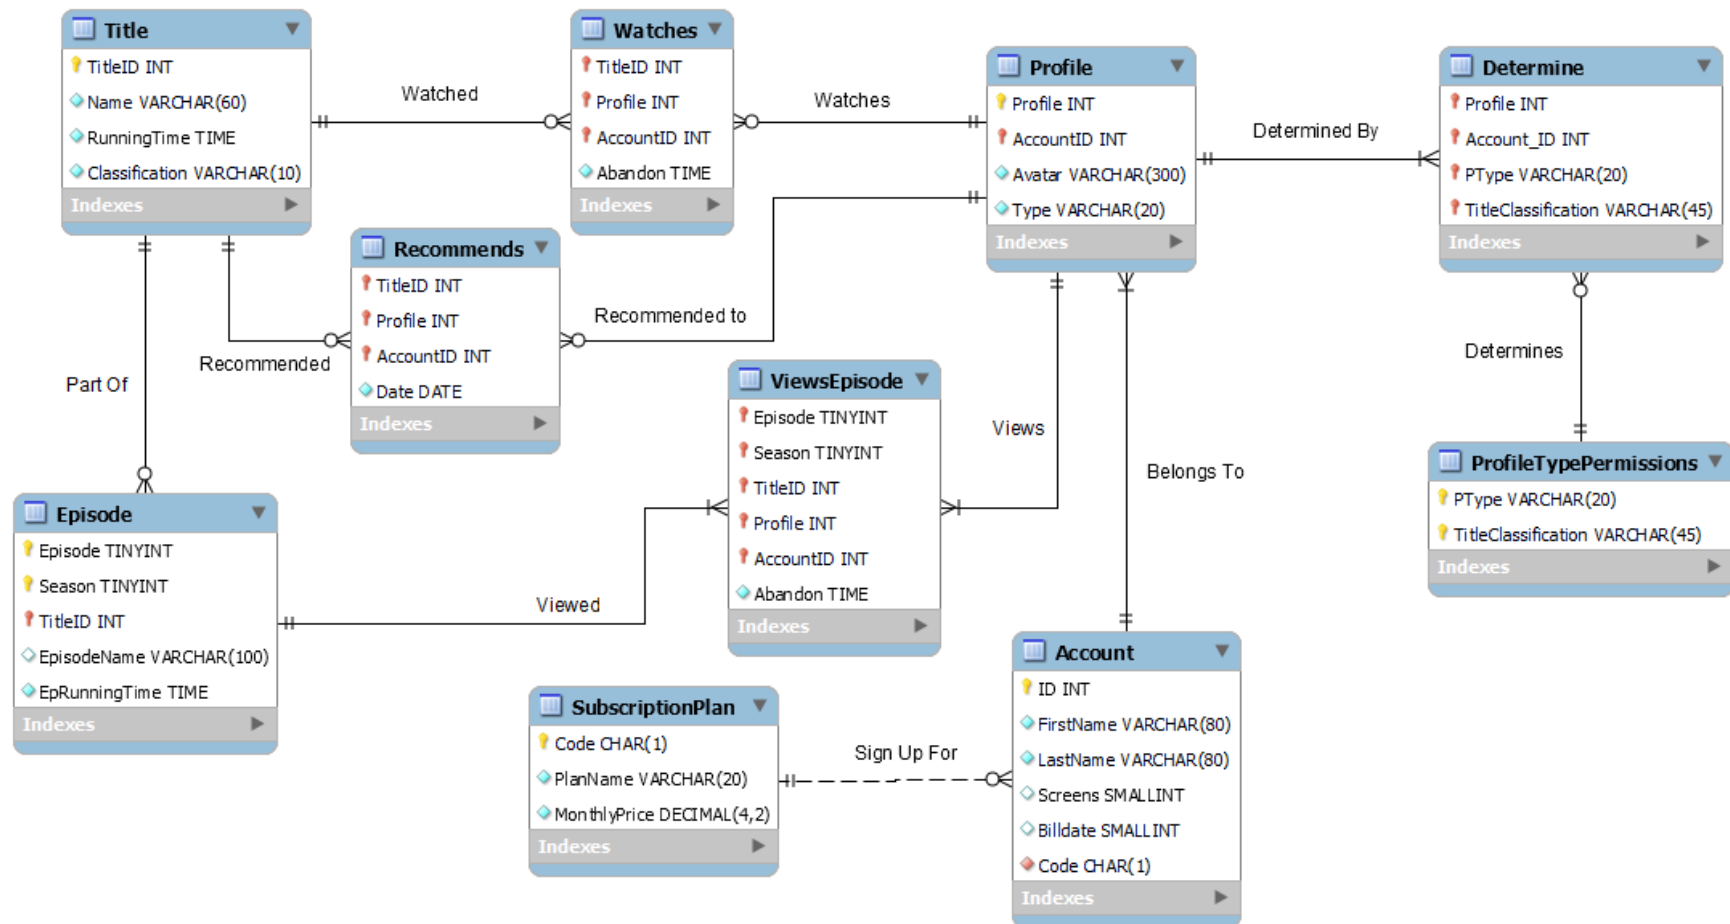
Each movie is an individual title. TV shows can have one or more seasons and up to 24 episodes in a season. Each title of a TV show (e.g. "Breaking Bad") will also have episodes title (e.g. "Ozymandias") as well as season number (5) and episode number (14) as well as its running time 47:35 (minutes and seconds) and classification ('MA').

Criterion Classics stores the length of each title/episode in minutes and seconds, as well as the elapsed time a particular profile stopped watching the title/episode. This is so users can 'continue watching' where they left off if they have to stop watching for any reason. If a user watches the title to the end, the abandon time is equal to the title/episode length time.

By collecting and storing this information, Criterion Classics can make recommendations to other users based on profile data analysis. Criterion Classics makes recommendations about rating appropriate titles to all profiles. Criterion needs to store the name of the title, the date it was suggested to a profile and the number of times (total count) a particular title has been recommended to a user. Using this information and the viewing information of a profile it can determine if its suggested viewing algorithm is working.

**Q1.** You are asked to model a **physical** Model of the Criterion Classics case study in **Crows Foot** notation. State any assumptions you have made.

# INFO90002 PRACTICE EXAM 1 SOLUTION



## Q2. SQL (15 marks)

*This case study is for the SQL model to help contextualise the Physical ER model*

### Lunch Rider

LunchRider is a new startup in the food delivery business. LunchRider allows people to order lunches from local food vendors and have them delivered by a delivery rider on a bike. When a customer opens the LunchRider app, it first presents a list of local vendors such as cafes, restaurants and snack bars. The customer clicks on a vendor. Then the app shows the meals available from that vendor.

The customer chooses which meals they want from that vendor: for example "2 chili burgers meals" and "1 mango smoothie". The phone sends this order to the LunchRider server, along with the customer's phone GPS coordinates - the meals will be delivered to this location. Customers can click to "like" a meal and the total number of likes (across all customers) is displayed beside each meal. When a customer's order is received, LunchRider broadcasts a work offer to riders who are near the customer. Riders see the offer pop up on their app and can press "accept" or "no thanks".

The chosen rider goes to the vendor, picks up the meals, and delivers them to the customer's location. We record at what time the rider delivers the order. Payment is automatically deducted from the customer's credit card and bank transactions are handled by the bank. LunchRiders provide their name, mobile phone number, and date of birth. Whenever a rider is on duty, the rider's app sends the GPS coordinates about once per minute, allowing us to keep track of each rider's location.

If someone wants to order a meal, they need to first register as a customer, giving their name, email address, mobile phone number and current credit card. Over time a customer may register more than one credit card. Vendors must register the name and address of their business, including GPS coordinates and a contact email address, and provide the name, price, description (max 1,000 characters) and photo of each meal that they want to sell.

After delivery, the app allows the customer to rate the rider's service, choosing from 1 star (worst), 2, 3, 4 or 5 stars (best). The app also allows the customer to add the rider to their "favourite" list. LunchRider uses this information to help choose riders for future work offers.

All locations are recorded as a pair of numbers representing latitude and longitude. Latitudes are between -90 and 90 degrees (south pole to north pole) while longitudes are between -180 and 180 degrees (west or east of the prime meridian). LunchRider uses a precision of 4 decimal places, which is about 11 metres at the equator (smaller in Melbourne). For example, the Doug McDonnell building at The University of Melbourne is at latitude -37.7989, longitude 144.9627.

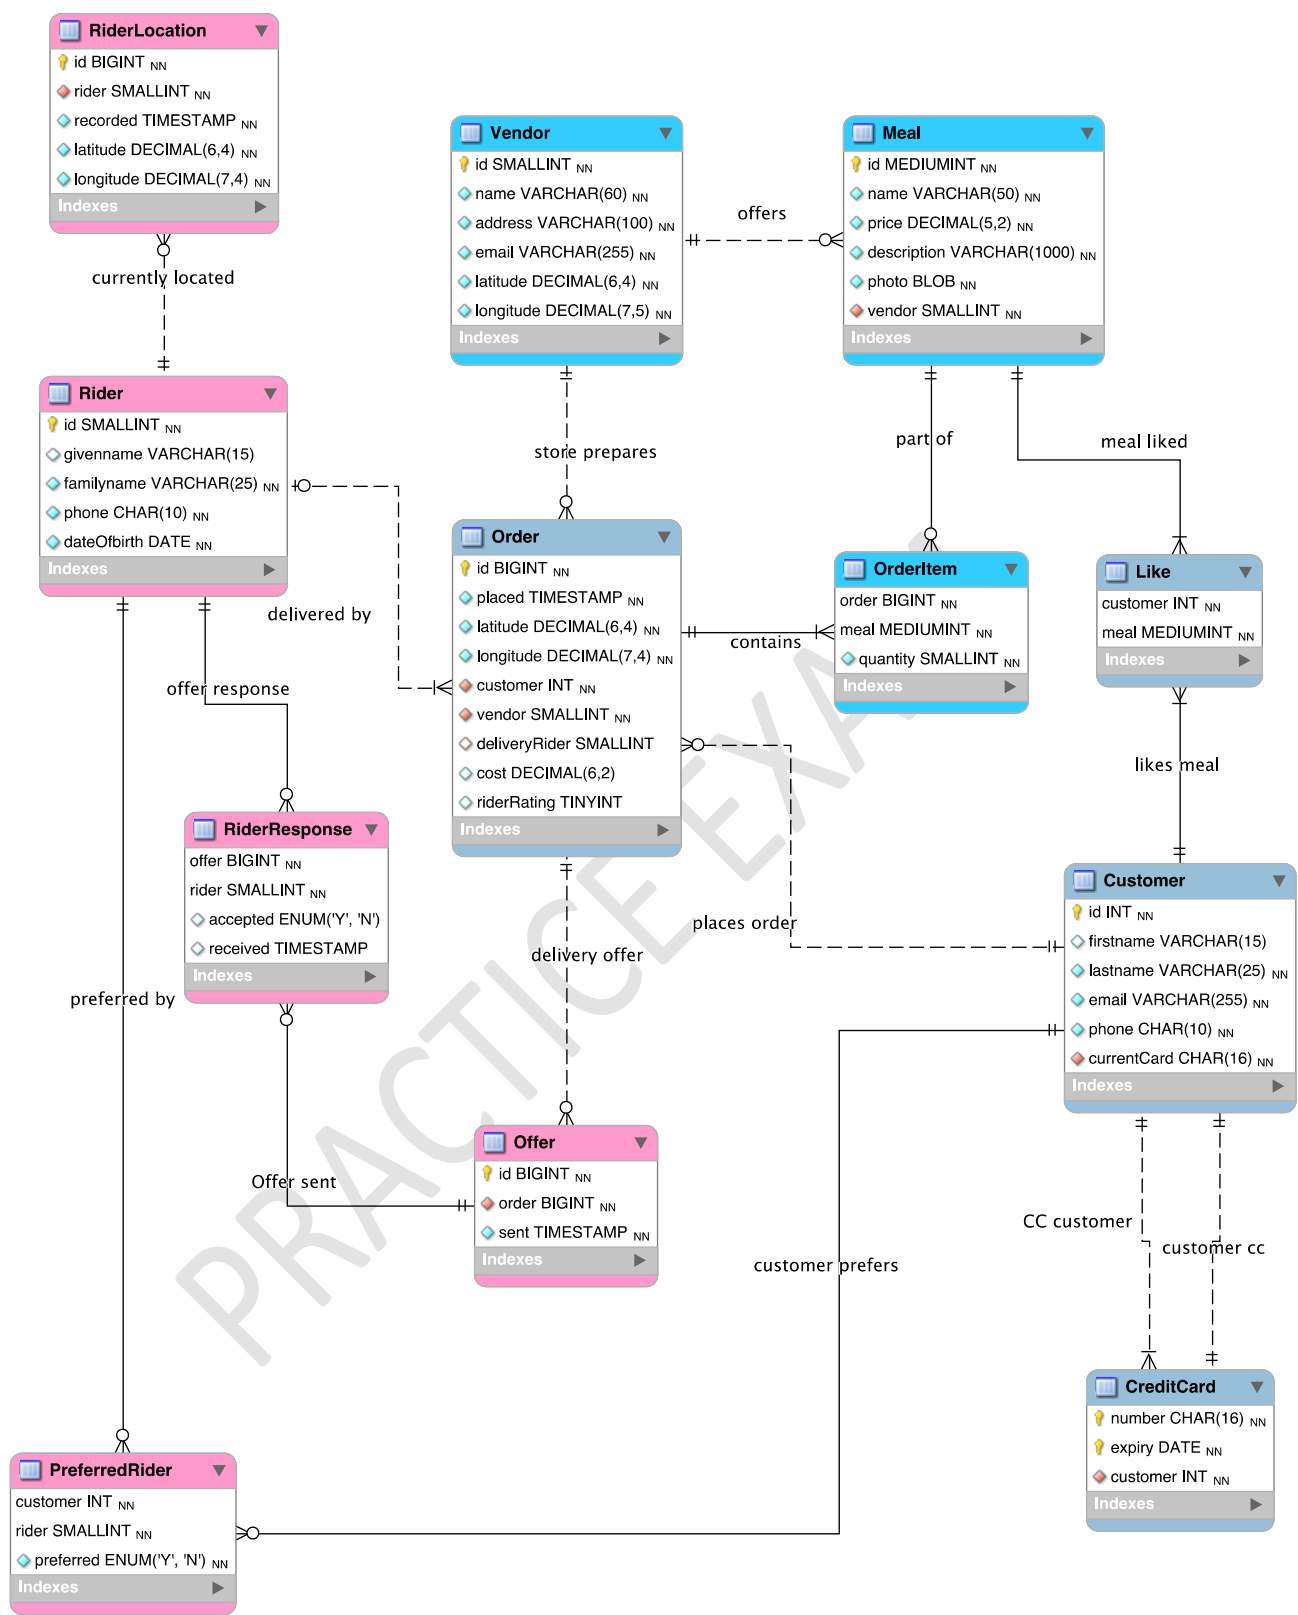


Figure 2: The Lunch Rider Entity Relationship Model

Questions 2A-2G require you to write one single SQL statement per question. Do **not** use views, temporary tables or inline views. Format code for ease of reading. Ensure user-friendly output by renaming columns where appropriate. Values of column fields are in italics e.g. *Tom Hardy*

For Example:

**Q.** List the name and salary of *Jane Grey*

```
SELECT firstname, lastname, salary
FROM employee
WHERE firstname = 'Jane' AND lastname = 'Grey';
```

**Q2A.** Write the SQL for the following query:

List all Vendor names who offer a meal where the price is more than \$60.00 and below \$80.00.

```
SELECT DISTINCT Vendor.Name
FROM Vendor
INNER JOIN Meal
ON Vendor.ID = Meal.vendor
WHERE Meal.price > 60 and Meal.price < 80;
```

**Q2B.** List the names of riders born in 1997 have accepted an offer in *August 2020*?

```
SELECT DISTINCT givenname, familyname
FROM rider INNER JOIN RiderResponse as RR
ON rider.id = RR.rider
WHERE RR.accepted = 'Y'
AND YEAR(rider.dateOfBirth) = 1997
AND MONTHNAME(RR.received) = 'August' AND YEAR(RR.received) = 2020;
```

**Q2C.** How many customers have ordered and liked the '*Betty Deluxe Burger Meal*' meal ordered from the Melbourne store with coordinates latitude -37.8156 and longitude 144.9636 ?

```
SELECT m.name, count(DISTINCT l.customer)
FROM Vendor v INNER JOIN Meal m INNER JOIN Like l INNER JOIN OrderItem oi
ON v.ID = m.vendor AND Meal.ID = l.Meal AND oi.Meal = m.ID
WHERE v.latitude = -37.8156 AND v.longitude = 144.9636
AND m.name = 'Betty Deluxe Burger Meal';
```

Note, GROUP BY m.name is not needed here because we are counting customers who ordered one specific meal.

**Q2D.** List the names of riders who are also customers of the Lunch Rider app (assume the rider uses the same phone number as they would as a customer).

```
SELECT firstname, lastname
FROM Customer
WHERE EXISTS
    (SELECT *
     FROM Rider
     Where Rider.phone = Customer.phone);
```

**Q2E.** List the vendor names who have at least one ordered meal never liked by any customer.

```
SELECT DISTINCT vendor.name
FROM VENDOR
INNER JOIN MEAL
INNER JOIN OrderItem
ON Vendor.ID = Meal.Vendor AND Meal.ID = OrderItem.meal
WHERE Meal.ID NOT IN
    (SELECT meal
     FROM Like);
```

-- (NOT EXISTS also possible)

**Q2F.** List the vendor names who have never had any ordered meal liked by a customer.

```
SELECT DISTINCT vendor.name
FROM VENDOR
INNER JOIN MEAL
INNER JOIN OrderItem
ON Vendor.ID = Meal.Vendor AND Meal.ID = OrderItem.meal
WHERE Vendor.ID NOT IN
    (SELECT DISTINCT Vendor.id
     FROM Vendor
     INNER JOIN Meal ON Vendor.ID = Meal.Vendor
     INNER JOIN Like ON Like.meal = Meal.id
    );
```

And a much better solution

```
SELECT Vendor.Name
FROM Vendor
INNER JOIN Meal ON Meal.vendor = Vendor.id
INNER JOIN OrderItem ON OrderItem.meal = Meal.id
LEFT OUTER JOIN 'Like' ON 'Like'.meal = Meal.id
GROUP BY Vendor.id, Vendor.Name
HAVING COUNT('Like'.meal) = 0;
```

**Q2G.** Name the riders who have received a rating of 4 or above from more than 500 different customers.

```
SELECT givenname, familyname
FROM Rider INNER JOIN Order
ON Rider.ID = Order.deliveryrider
WHERE Order.riderrating >= 4
GROUP BY Rider.ID
HAVING COUNT(DISTINCT(Order.customer)) > 500;
```



### Q3. Normalisation (15 marks)

The table below is part of the medical records for a Veterinarian Clinic

PRACTICE (AnimalID, AnimalName, AnimalType, owner\_ID, Owner, Phone, ConsultDateTime, procID, procDescription)

Animal ID	Animal name	Animal Type	Owner_ID	Owner	Phone	ConsultDateTime	ProcID	Description
317	Ralph	Dog	10	Julie Sumner	0409 673-888	13-Oct-22 9:00	101	Annual Checkup
317	Ralph	Dog	10	Julie Sumner	0409 673-888	27-Apr-22 11:15	115	Teeth Clean
317	Ralph	Dog	10	Julie Sumner	0409 673-888	14-Oct-23 9:30	119	3 month Checkup
398	Zeno	Canary	23	Tony Rijks	0408 322-444	21-Jul-23 14:30	105	Parasite treatment
398	Zeno	Canary	23	Tony Rijks	0408 322-444	14-Oct-23 16:00	119	3 month Checkup
441	Panda	Short haired cat	47	Helene Hanff	0419 121-212	24-Apr-23 9:00	715	Initial Consultation
441	Panda	Short haired cat	47	Helene Hanff	0419 121-212	27-Apr-23 9:30	115	Teeth Clean
398	Zeno	Canary	23	Tony Rijks	0408 322-444	1-Mar-24 9:00	001	6 month Checkup

The combination of AnimalID and procID is the candidate key for the relation. The following functional dependencies hold:

- AnimalID → animalName, animalType
- owner\_id → owner, phone
- procid → description
- animalID, procID → consultDateTime

Assume more than one procedure can be performed during the consultation.

**Q3.** Normalise the data to third normal form (3NF). Be sure to identify all anomalies, and show each stage (1NF, 2NF, 3NF).

Key: **Underline** primary key *ITALIC* foreign key **Underline + ITALIC** primary foreign key

### 1NF All atomic cells

However Insert Update Delete Anomalies exist

- DELETE Zeno the Canary we lose the 6-month checkup procedure (001)
- UPDATE Ralph's owner – multiple updates
- INSERT Have to know the consultation date/time before we can add owner and pet

### 1NF

PRACTICE(animalID, animalName, animalType, owner\_ID, owner, phone, (procID, ProcDescription, consultDateTime)) – identify repeating groups

Remove repeating groups (note renaming entities makes better entity names):

Animal (animalID, animalName, animalType, owner\_ID, owner, phone)

Consultation (procID, ProcDescription, consultDateTime, animalID)

### 2NF

To be in 2NF - 1NF and No partial functional dependencies

Partial functional dependency exists in Consultation

ProcID → ProcDescription

Animal (animalID, animalName, animalType, owner\_ID, owner, phone)

Procedure (procID, ProcDescription)

Consultation (procID, animalID, consultDateTime)

Anomalies

If the owner changes we have multiple updates

### 3NF

To be in 3NF - 2NF and no transitive functional dependencies

Transitive functional dependencies exist

Owner\_ID --> owner, phone - not dependent on animalID of Animal

Procedure (procID, ProcDescription)

Consultation (procID, consultDateTime, animalID)

Animal (animalID, animalName, animalType, owner\_ID)

Owner (owner\_ID, owner, phone)

PRACTICE EXAM

**Q4. NoSQL****(4 marks)**

**Q4.** There are trade-offs between the principles of ACID and BASE. Discuss the trade-off between availability and consistency for relational and NoSQL databases. Illustrate your answer using either the example of Facebook or the example of Twitter.

**( 4 marks)**

NoSQL databases need to be available all the time. Facebook (or Twitter) needs to ensure that users can access and post to its site and app all the time, even if the data is not consistent. For example, users expect to always see the number of likes on a post (or retweet count on a tweet); they do not mind if that number is slightly outdated.

Relational databases sacrifice availability for consistency. This means the database needs to be consistent all the time for all the users. If there is potential for inconsistencies the system may shutdown, this would cause severe disruption to Twitter/Facebook which rely on being always available.

**Q5. Applications****(4 marks)**

**Q5.** There are problems with giving end users an SQL interface to access a database. Describe **two** (2) distinct problems, and for each, how providing application software to users solves the problem.

**(2+2=4 marks)**

Providing users with an interface to a database presents two distinct problems. The first is users would be able to make indiscriminate changes including errors, and would all need to be trained in the SQL language

The second is logically units of work. By providing an application interface you can structure SQL into transactions such as purchasing goods online, transferring funds, or adding student results for subjects.

**Q6. Distributed Databases****(4 marks)**

**Q6.** Describe how synchronous updates work and the advantages and disadvantages of this approach.

**(4 marks)**

Synchronous updates apply to Distributed Replicated databases. When a change is made on one node/server in the distributed database the change is immediately sent to all other nodes in the distributed database. All users on all nodes see the same current value for all data immediately – therefore providing excellent data integrity

The problem with synchronous replication is that any break in the network would cause an outage, as data will not be consistent. Synchronous replication increases network traffic and may result in slower performance.

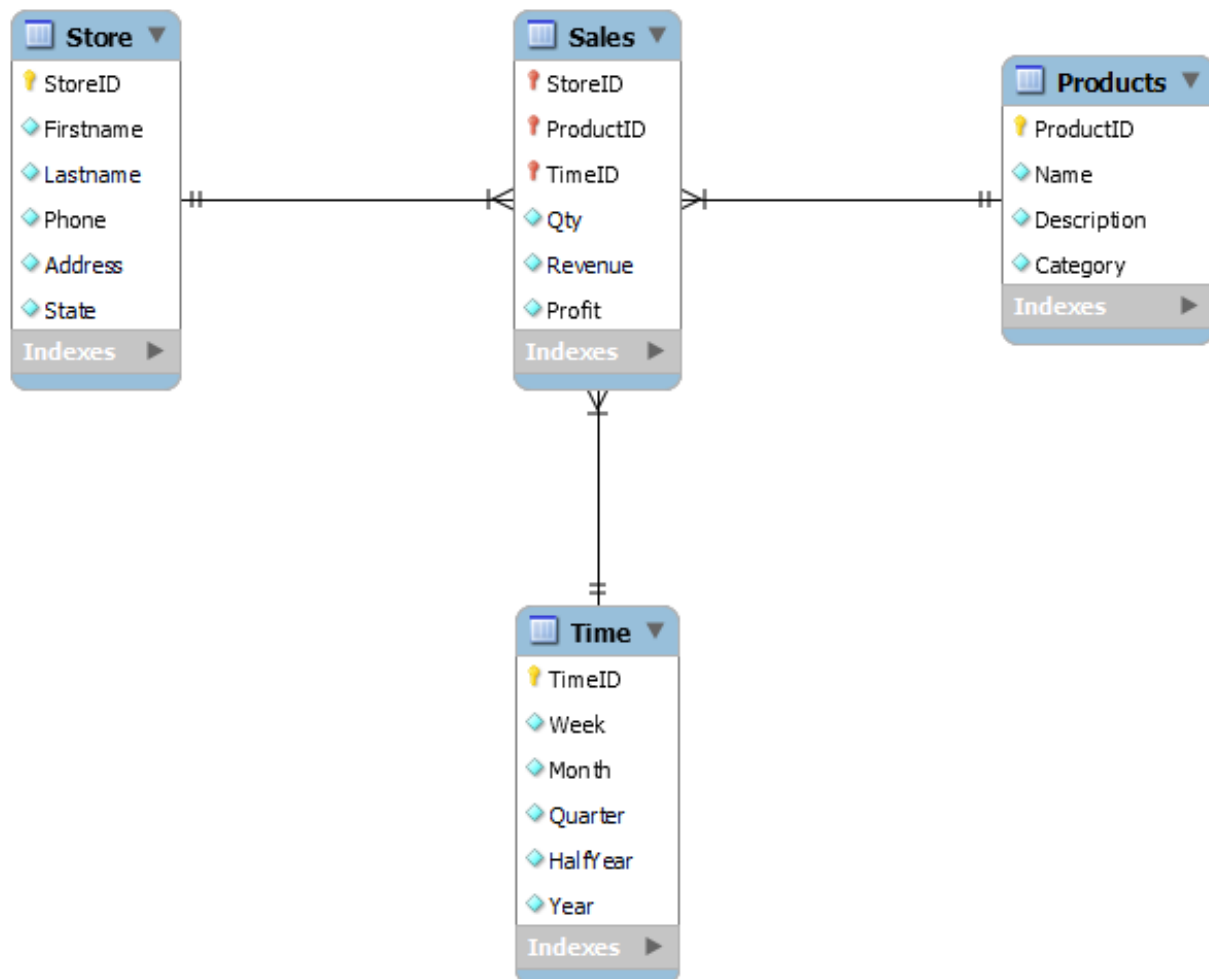
## Q7. Data Warehouses

(14 Marks)

Sennheiser Australia wants to track information about the selling of audio equipment (microphones, headphones, turntables) via physical retail stores located throughout Australia. You need to design a data warehouse to report information about sales of audio equipment over time. You need to store the retail store (store ID, Manager's first name, Manager's last name, phone number, address, state), and product (product-id, name, description, category, retail price, wholesale price). The sales managers want to find the number of items sold (e.g. turntables), the revenue and profit. The information needs to be accessible by the retail store, its state, product name (e.g. "Sen-300-II") and for different times (week, month, quarter, half year and year).

**Q7A.** Draw a *star schema* to support the design of this data warehouse, showing the attributes in each table. Be sure to denote PK, FK and PFK.

(8 marks)



**Q7B.** Briefly explain what OLTP and OLAP databases are and demonstrate the difference between these two types of databases.

(6 marks)

OLTP are online transactional processing databases designed to collect the data associated with day to day operations of the business. The database design is normalised and the data is volatile in nature and subject to changes. Data is both read and written.

OLAP are online analytical processing databases. They are designed to analyse historic data that is largely static and in a finalised form. The majority of transactions are read only - and long running, and designed for management and strategic - not operational decision making. OLAP databases are usually focused on a single theme (e.g Sales).

The difference between the two systems are listed below

Characteristic	OLTP	OLAP
Data	Volatile	Finalised
Transactions	Read and Write; Many short transactions	Read only Long running
Database Use	Operational	Management and Strategic
Users	Many users for day to day operations	Managers to make strategic decisions

**Q8. Ethics & Security (4 marks)**

**Q8.** What are the advantages and disadvantages of a logical backup when compared with a physical backup of a database?

(4 marks)

The advantage of a logical backup is that it can occur while the database is still running - and therefore not interrupt database availability. The backup is in a series of SQL statements than can be ported to any operating system.

The disadvantages of a logical backup is that it may impact database performance as it is resource intensive. Logical backups are larger in size than physical backups and usually take longer to complete.

END OF EXAM  
GOOD LUCK!