



THE UNIVERSITY OF
MELBOURNE

SQL Part 1

Database Systems & Information Modelling
INFO90002.

Week 4 - SQL

Dr Tanya Linden
Dr Renata Borovica-Gajic
David Eccles





Changing your INFO90002db password

Student passwords are the username and then the year

E.g. [jwu3](#) username would be [jwu3_2024](#) password

To change your password:

- Connect to the INFO90002db server (you will need the VPN outside lab computers)

SYNTAX

```
SET PASSWORD = 'newpassw0rd';
```



Installing the Labs Script

rename 9999 → to studentID

BYOD (both workbench and MySQL Server are on the same machine)

The install script for the labs has two versions

Use labs2018byod.sql if you are using your own server, on your own computer

Use labs2018eng.sql if you are using the server on info90002db.eng.unimelb.edu.au

Note 1: students do not have permissions to run CREATE DATABASE command on the engineering server

Note 2: do not install assignment 2 scripts before reading instructions (you will regret it ;-)



What is SQL

- SQL – or "sequel" is a language used in relational databases
- ALL DBMS support CRUD
 - Create, Read, Update, Delete commands
- SQL supports CRUD
 - CREATE, SELECT, INSERT, UPDATE, DELETE, DRQP commands
- Other info
 - You can see the 2011 standard of SQL at
 - http://www.jtc1sc32.org/doc/N2151-2200/32N2153T-text_for_ballot-FDIS_9075-1.pdf
 - Wikipedia has several sections on SQL (good for generic syntax)
 - http://en.wikipedia.org/wiki/Category:SQL_keywords
- W3Schools has a good self-paced tutorial and examples
 - <https://www.w3schools.com/sql/default.asp>



SQL Language

- Provides the following capabilities:
 - Data Definition Language (DDL)
 - To define and set up the database
 - CREATE, ALTER, DROP
 - Data Manipulation Language (DML)
 - To maintain and use the database
 - SELECT, INSERT, DELETE, UPDATE
 - Data Control Language (DCL) *not learning in this subject*
 - To control access to the database
 - GRANT, REVOKE
 - Other Commands
 - Administer the database
 - Transaction Control
 - START TRANSACTION
 - BEGIN, END



How We Use SQL

- In **Implementation** of the database
 - Take the tables we design in physical design
 - Implement these tables in the database using create commands
- In **Use** of the database
 - Use SELECT commands to read the data from the tables, link the tables together, etc.
 - Use ALTER, DROP commands to update the database
 - Use **INSERT, UPDATE, DELETE** commands to change data in the database



CREATE, INSERT, NULL

SQL DDL and DML

SQL in Development Process

1.

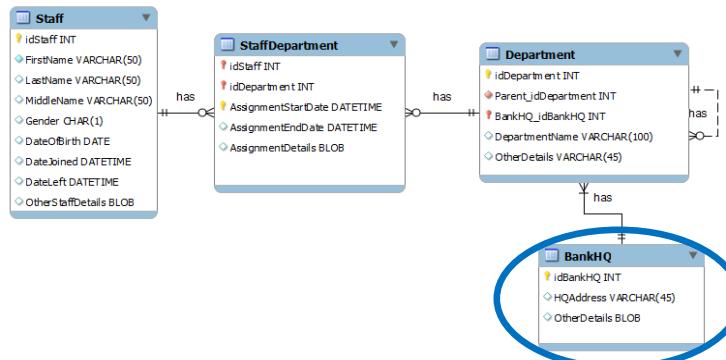
```
CREATE TABLE BankHQ (
    BankHQID          INT          AUTO_INCREMENT,
    HQAddress         VARCHAR(120) NOT NULL,
    OtherHQDetails   VARCHAR(100),
    PRIMARY KEY(BankHQID)
);
```
2.

```
INSERT INTO BankHQ VALUES
(1, "23 Charles St Peterson North 2022", "Main Branch");
INSERT INTO BankHQ VALUES
(2, "213 Jones Rd Parkville North 2122", "Sub Branch");
```

You have to input every column
3.

```
SELECT *
FROM BankHQ;
```

BankHQID	HQAddress	OtherHQDetails
1	23 Charles St Peterson North 2022	Main Branch
2	213 Jones Rd Parkville North 2122	Sub Branch



CREATE Table and Viewing Table Definition

Customer

- CustomerID NN
- CustFirstName
- CustMiddleName
- CustLastName NN
- BusinessName
- CustType NN

Indexes

```
CREATE TABLE Customer (
  CustomerID INT AUTO_INCREMENT,
  CustFirstName VARCHAR(12),
  CustMiddleName VARCHAR(14),
  CustLastName VARCHAR(20),
  BusinessName VARCHAR(100),
  CustType ENUM("Personal","Company")
  PRIMARY KEY(CustomerID)
);
```

AUTO_INCREMENT,

NOT NULL,

NOT NULL,

To check what you created use one of the following:

DESCRIBE Customer;

see the describe

DESC Customer;

SHOW COLUMNS FROM Customer;

the result will be the same.

Field	Type	Null	Key	Default	Extra
CustomerID	int	NO	PRI	NULL	auto_increment
CustFirstName	varchar(12)	YES		NULL	
CustMiddleName	varchar(14)	YES		NULL	
CustLastName	varchar(20)	NO		NULL	
BusinessName	varchar(100)	YES		NULL	
CustType	enum('Personal','Company')	NO		NULL	



Insert Data

```
INSERT INTO Customer (CustFirstName, CustLastName, CustType)  
VALUES ("Peter", "Smith", "Personal");
```

Specifies which columns will be entered; ID will be auto-incremented

```
INSERT INTO Customer  
VALUES (DEFAULT, "James", NULL, "Jones", "JJ Enterprises", "Company");
```

```
INSERT INTO Customer  
VALUES (DEFAULT, "", NULL, "Smythe", "", "Company");
```

No column specification means ALL columns need to be entered

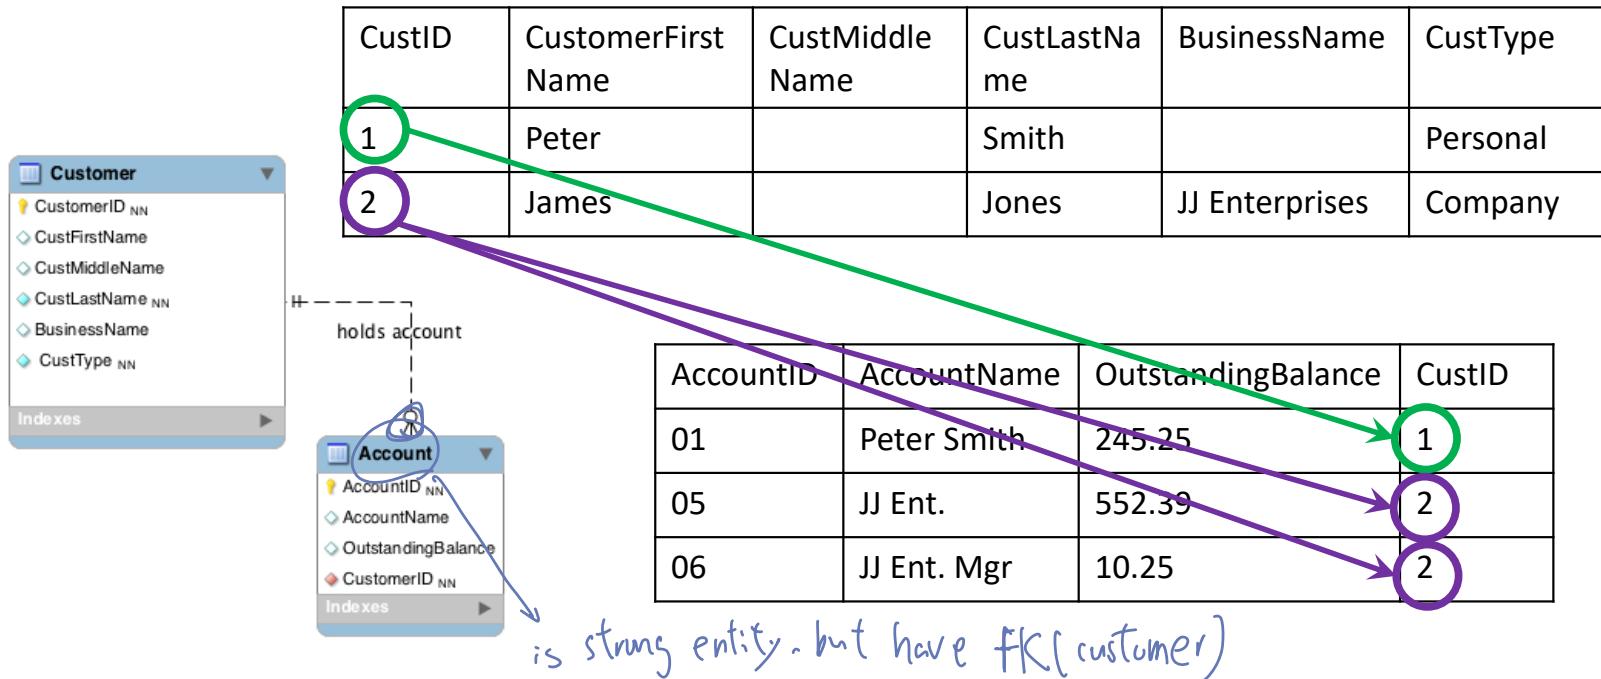
Customer

CustID	CustomerFirstName	CustMiddleName	CustLastName	BusinessName	CustType
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3		NULL	Smythe		Company

What does NULL mean? Null Island: The Busiest Place That Doesn't Exist by MinuteEarth
<https://www.youtube.com/watch?v=bjvlpI-1w84>

Foreign keys: Review

- We looked at Customer
 - A customer can have a number of Accounts
 - The tables get linked through a foreign key





CREATE Statement (with FK)

In SQL, the `ON DELETE RESTRICT` constraint is used to prevent the deletion of a row from a parent table if there are corresponding rows in a child table that reference it. This constraint ensures referential integrity by enforcing that the parent row cannot be deleted as long as there are dependent rows in the child table.

```
CREATE TABLE Account (
    AccountID          INT          AUTO_INCREMENT,
    AccountName        VARCHAR(12),
    OutstandingBalance DECIMAL(10,2) NOT NULL,
    CustomerID         INT          NOT NULL,
    PRIMARY KEY (AccountID),
    FOREIGN KEY (CustomerID) REFERENCES Customer(CustomerID)
        ON DELETE RESTRICT
        ON UPDATE CASCADE
)
```

primary key in customer



SELECT

SQL DML

The SELECT Statement

A cut down version of the SELECT statement – MySQL

SELECT [ALL | DISTINCT] *select_expr* [, *select_expr* ...]

- List the columns (and expressions) that are returned from the query

[FROM *table_references*]

- Indicate the table(s) or view(s) from where the data is obtained

[WHERE *where_condition*]

- Indicate the conditions on whether a particular row will be in the result

[GROUP BY {*col_name* | *expr*}, ...]

- Indicate categorisation of results

[HAVING *where_condition*]

- Indicate the conditions under which a particular category (group) is included in the result

[ORDER BY {*col_name* | *expr* | *position*} [ASC | DESC], ...]

- Sort the result based on the criteria

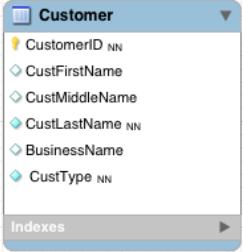
[LIMIT {[*offset*,] *row_count* | *row_count* OFFSET *offset*}]

- Limit which rows are returned by their return order (i.e. 5 rows, 5 rows from row 2)



Order is important! E.g. **HAVING cannot go before GROUP BY or WHERE**

SELECT Examples



SELECT * FROM Customer;
= Extract all data about customers

SQL **SELECT ***
FROM Customer

RESULT
SET

	CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType
▶	1	Peter	NULL	Smith	NULL	Personal
	2	James	NULL	Jones	JJ Enterprises	Company
	3	Akin	NULL	Smithies	Bay Wart	Company
	4	Julie	Anne	Smythe	Konks	Company
	5	Jen	NULL	Smart	BRU	Company
	6	Lim	NULL	Lam	NULL	Personal
	7	Kim	NULL	Unila	Saps	Company
	8	James	Jay	Jones	JJ's	Company
	9	Keith	NULL	Samson	NULL	Personal

SELECT Examples: Projection

Projection means specifying which attributes to show in the results set.

Customer	
CustomerID	NN
CustFirstName	
CustMiddleName	
CustLastName	NN
BusinessName	
CustType	NN
Indexes	

In SQL:

```
SELECT CustLastName  
FROM Customer;
```

end of instruction

Result set

in the instruction
separate by (omc)

CustLastName
Smith
Jones
Smithies
Smythe
Smart
Lam
Unila
Jones
Samson



SELECT Examples: Selection

In SQL:

```
SELECT CustLastName  
FROM Customer  
WHERE CustLastName = "Smith";
```

Selection means specifying conditions that records in the output must meet.

Result

CustLastName
▶ Smith

SELECT Examples: LIKE clause

In addition to arithmetic expressions, string conditions are specified with the LIKE clause

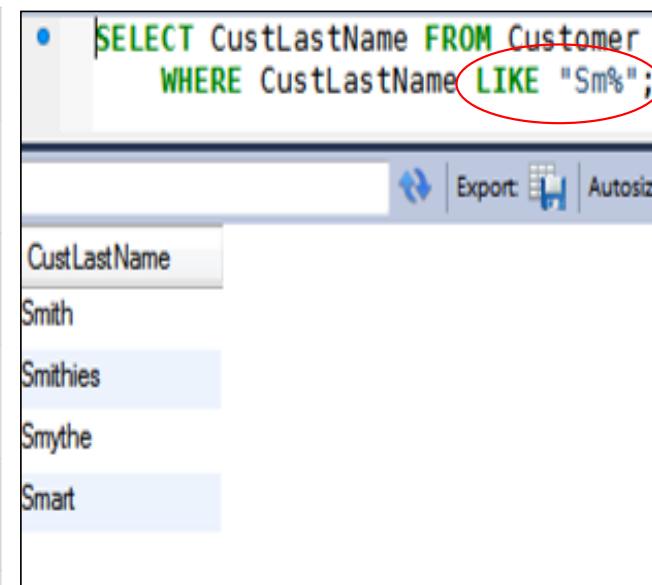
LIKE "REG_EXP"

% Represents zero, one, or multiple characters
_ Represents a single character

Examples:

WHERE CustLastName LIKE 'a%'	Finds any values that start with "a"
WHERE CustLastName LIKE '%a'	Finds any values that end with "a"
WHERE CustLastName LIKE '%or%'	Finds any values that have "or" in any position
WHERE CustLastName LIKE '_r%'	Finds any values that have "r" in the second position
WHERE CustLastName LIKE 'a_%_%'	Finds any values that start with "a" and are at least 3 characters in length
WHERE CustLastName LIKE 'a%o'	Finds any values that start with "a" and end with "o"

SQL:



The screenshot shows a SQL query being run against a 'Customer' table. The query is:

```
SELECT CustLastName FROM Customer  
WHERE CustLastName LIKE "Sm%";
```

The result set displays several last names:

CustLastName
Smith
Smithies
Smythe
Smart

Aggregate Functions

Aggregate functions operate on the (sub)set of values in a column of a relation (table) and return a single value

- COUNT()
 - Number of values
 - SUM()
 - Sum of values
 - AVG()
 - Average value
 - MIN()
 - Minimum value
 - MAX()
 - Maximum value
-
- 
- N.B. All of these except for COUNT() ignore null values and return null if all values are null. COUNT() counts the rows not the values and thus even if the value is NULL, it is still counted
 - COUNT(Column) versus COUNT(*)
 - List of MySQL aggregate functions
 - <https://dev.mysql.com/doc/refman/8.0/en/aggregate-functions.html>



Aggregate Examples: COUNT and AVG

- | | |
|---------|---------------------------------|
| COUNT() | - returns the number of records |
| AVG() | - average of the values |

Examples:

```
SELECT COUNT(CustomerID)  
FROM Customer;
```

= How many customers do we have
(cardinality)

```
SELECT AVG(OutstandingBalance)  
FROM Account;
```

= What is the average balance of
ALL ACCOUNTS

```
SELECT AVG(OutstandingBalance)  
FROM Account  
WHERE CustomerID=1;
```

= What is the average balance of
Accounts of Customer with ID 1

where use in attribute

```
SELECT AVG(OutstandingBalance)  
FROM Account  
GROUP BY CustomerID;
```

= What is the average balance
PER CUSTOMER



GROUP BY / HAVING

Group by groups all records together over a set of attributes

Frequently used with aggregate functions

Example:

*What is the average balance **PER CUSTOMER**?*

```
SELECT AVG(OutstandingBalance)  
FROM Account  
GROUP BY CustomerID;
```

The only way to put a selection condition over a group by statement is by using **having** clause

Example:

What is the exact average balance per customer for customers whose average balance is under 10000?

```
SELECT AVG(OutstandingBalance)  
FROM Account  
GROUP BY CustomerID  
HAVING AVG(OutstandingBalance) < 10000
```

for aggregate function you can not use where.



Changing Column Heading in Output - Alias

We can rename the column name in the output by using the AS clause

If it contains a gap, it must be in straight double quotes

```
SELECT custtype AS "Customer Type", COUNT(customerid) AS CUST_TOTAL  
FROM customer  
GROUP BY custtype;
```

more than two words
use "

```
SELECT custtype, COUNT(customerid)  
FROM customer  
GROUP BY custtype;
```

custtype	COUNT(customerid)
Personal	1
Company	2

alias to userNcmpl

Customer Type	CUST_TOTAL
Personal	1
Company	2



ORDER BY

- Orders records by particular column(s)

ORDER BY colName ASC/DESC (ASC is default)

```
SELECT firstname, lastname  
FROM employee  
ORDER BY lastname;
```

firstname	lastname
Todd	Beamer
Nancy	Cartwright
Pat	Clarkson
Sarah	Fergusson
Paul	Innit
Andrew	Jackson
Ned	Kelly
James	Mason
Sophie	Monk
Gigi	Montez
Alice	Munro
Brier	Patch
Sanjay	Patel
Rita	Skeeter
Maggie	Smith
Clare	Underwood
Mark	Zhang

by default ascend

```
SELECT firstname, lastname, departmentid  
FROM employee  
ORDER BY departmentid DESC, lastname ASC;
```

firstname	lastname	departmentid
Andrew	Jackson	11
Ned	Kelly	11
Clare	Underwood	11
Sophie	Monk	10
Sarah	Fergusson	9
Brier	Patch	9
Todd	Beamer	8
Nancy	Cartwright	8
Mark	Zhang	7
Sanjay	Patel	6
Pat	Clarkson	5
Paul	Innit	4
James	Mason	4
Gigi	Montez	3
Maggie	Smith	3
Rita	Skeeter	2
Alice	Munro	1

suit your result by particular order

first order
second order
optional to write

LIMIT and OFFSET

LIMIT N - limits the output size

OFFSET N - skips first N records

```
SELECT firstname, lastname, departmentid
FROM employee
ORDER BY departmentid DESC, lastname ASC
LIMIT 6;
```

firstname	lastname	departmentid
Andrew	Jackson	11
Ned	Kelly	11
Clare	Underwood	11
Sophie	Monk	10
Sarah	Fergusson	9
Brier	Patch	9

LIMIT 6

firstname	lastname	departmentid
▶ Andrew	Jackson	11
Ned	Kelly	11
Clare	Underwood	11
Sophie	Monk	10
Sarah	Fergusson	9
Brier	Patch	9
Todd	Beamer	8
Nancy	Cartwright	8
Mark	Zhang	7
Sanjay	Patel	6
Pat	Clarkson	5
Paul	Innit	4
James	Mason	4
Gigi	Montez	3
Maggie	Smith	3
Rita	Skeeter	2
Alice	Munro	1

```
SELECT firstname, lastname, departmentid
FROM employee
ORDER BY departmentid DESC, lastname ASC
LIMIT 6 OFFSET 3;
```

firstname	lastname	departmentid
▶ Sophie	Monk	10
Sarah	Fergusson	9
Brier	Patch	9
Todd	Beamer	8
Nancy	Cartwright	8
Mark	Zhang	7

LIMIT 6 OFFSET 3

be careful if there are two highest
you have to use max() function.



Joining Tables

Cross Product,
INNER JOIN,
NATURAL JOIN,
RIGHT OUTER JOIN,
LEFT OUTER JOIN



Table name qualification and aliases

A simple SQL query using the SELECT statement

```
SELECT CustLastName, CustType  
FROM Customer;
```

Whenever we use a column name in a Query, we can qualify (prefix) with the appropriate table name

```
SELECT Customer.CustLastName, Customer.CustType  
FROM Customer;
```

We can use **an alias for table name**

```
SELECT c.CustLastName, c.CustType  
FROM Customer c;
```

The alias only applies to the current SQL statement (not remembered)

Once an alias is used within an SQL statement, you can **not** refer to the original tablename within that SQL statement

The usefulness of this feature becomes apparent when SQL statements refer to columns from two or more tables



Joining tables together

SELECT * FROM Table1, Table2; - the result set is a **cross product** or **Cartesian Product**

SELECT * FROM Customer, Account;

Customer, Account									
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	1	Peter Smith	245.25	1
3	Akin	NULL	Smithies	Bay Wart	Company	1	Peter Smith	245.25	1
1	Peter	NULL	Smith	NULL	Personal	2	JJ Ent.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	2
3	Akin	NULL	Smithies	Bay Wart	Company	2	JJ Ent.	552.39	2
1	Peter	NULL	Smith	NULL	Personal	3	JJ Ent. Mgr	10.25	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	2
3	Akin	NULL	Smithies	Bay Wart	Company	3	JJ Ent. Mgr	10.25	2

Every row in the first table has been matched with every row in the second table
Not very useful...

Reason: an **SQL Select** statement does **not** 'know' how two tables are **related** (it doesn't consider any existing FK – PK constraints)



Joins: INNER JOIN

An **INNER JOIN** returns a result set that contains only data that satisfies a **Foreign Key – Primary Key condition**

Syntax:

```
SELECT <column-names>
FROM   <table-name1>
INNER JOIN <table-name2>
ON     <join-condition>
```

The **<join-condition>** is normally in the format

<foreign-key column-name> = <primary-key column name>

```
SELECT      c.CustomerFirstName, c.CustomerLastName, a.OutstandingBalance
FROM        Customer c
INNER JOIN  Account a
ON          c.CustomerID = a.CustomerID
```

customer with the account won't be
shown in the joint table.

Joins: Different Types

Inner/Equi join:

- Joins the tables over keys

`SELECT * FROM Customer INNER JOIN Account
ON Customer.CustomerID = Account.CustomerID; CONDITION`

CustomerID CustFirstName CustMiddleName CustLastName BusinessName CustType AccountID AccountName OutstandingBalance CustomerID									
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	2

Natural Join:

- Joins the tables over keys. The condition does not have to be specified (NATURAL JOIN does it automatically), but key attributes have to have the *same name*.

natural join works
slower

`SELECT * FROM Customer NATURAL JOIN Account;`

CustomerID CustFirstName CustMiddleName CustLastName BusinessName CustType AccountID AccountName OutstandingBalance									
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	



Avoid **old** style Joins

Back in the **dark ages** (pre 2000), some DBMS products **did NOT use/have** the **INNER JOIN** keyword.

Instead the JOIN was expressed as part of the **where clause** within the Select statement

Many **old-time** developers, **old-time** web sites, **old-time** books / authors still use this old-style in their code and in their examples

```
SELECT c.CustomerFirstName, c.CustomerLastName, a.OutstandingBalance  
FROM Customer c, Account a  
WHERE c.CustomerID = a.CustomerID
```





Joins: Different Types

inner join:
both have to be satisfied

OUTER JOIN:

- Joins the tables over keys
- Can be *LEFT* or *RIGHT* (see difference below)
- Includes records that **don't match** the join from the other table

(left
customer is major

right
account is major)

SELECT * FROM Customer LEFT OUTER JOIN Account ON Customer.CustomerID = Account.CustomerID;									
Autosize: <input type="button" value="IA"/>									
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	HULL	Smith	HULL	Personal	1	Peter Smith	245.25	1
2	James	HULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	2
2	James	HULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	2
3	Akin	HULL	Smithies	Bay Wart	Company	HULL	HULL	HULL	HULL

Includes Customers with associated accounts as well as Customers who do not have Account

Vs

Customers with associated accounts as well as Accounts without Customers

SELECT * FROM Customer RIGHT OUTER JOIN Account ON Customer.CustomerID = Account.CustomerID;									
Autosize: <input type="button" value="IA"/>									
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	HULL	Smith	HULL	Personal	1	Peter Smith	245.25	1
2	James	HULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	2
2	James	HULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	2



What's examinable

- DDL
- DML
- SELECT



THE UNIVERSITY OF
MELBOURNE

Thank you